

Ridge and Lasso Regression

Advertising dataset

In [19]:

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LinearRegression
7 from sklearn.linear_model import Ridge, RidgeCV, Lasso
8 from sklearn.preprocessing import StandardScaler
```

In [20]:

```
1 data=pd.read_csv(r"C:\Users\91949\Downloads\Advertising.csv")
2 data
```

Out[20]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [21]:

```
1 data.head()
```

Out[21]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [22]:


```
1 data.tail()
```

Out[22]:


	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

```
In [23]: 1 plt.figure(figsize = (10, 10))
        2 sns.heatmap(data.corr(), annot = True)
```

Out[23]: <Axes: >



```
In [24]: 1 data.drop(columns = ["Radio", "Newspaper"], inplace = True)
        2 #pairplot
        3 sns.pairplot(data)
        4 data.Sales = np.log(data.Sales)
```



```
In [25]: 1 features = data.columns[0:2]
2 target = data.columns[-1]
3 #X and y values
4 X = data[features].values
5 y = data[target].values
6 #split
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
8 print("The dimension of X_train is {}".format(X_train.shape))
9 print("The dimension of X_test is {}".format(X_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 X_train = scaler.fit_transform(X_train)
13 X_test = scaler.transform(X_test)
```

The dimension of X_train is (140, 2)
The dimension of X_test is (60, 2)

```
In [26]: 1 #Model
2 lr = LinearRegression()
3 #Fit model
4 lr.fit(X_train, y_train)
5 #predict
6 #prediction = lr.predict(X_test)
7 #actual
8 actual = y_test
9 train_score_lr = lr.score(X_train, y_train)
10 test_score_lr = lr.score(X_test, y_test)
11 print("\nLinear Regression Model:\n")
12 print("The train score for lr model is {}".format(train_score_lr))
13 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:


The train score for lr model is 1.0
The test score for lr model is 1.0

```
In [27]: 1 #Ridge Regression Model
2 ridgeReg = Ridge(alpha=10)
3 ridgeReg.fit(X_train,y_train)
4 #train and test scorefor ridge regression
5 train_score_ridge = ridgeReg.score(X_train, y_train)
6 test_score_ridge = ridgeReg.score(X_test, y_test)
7 print("\nRidge Model:\n")
8 print("The train score for ridge model is {}".format(train_score_ridge))
9 print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.990287139194161
The test score for ridge model is 0.9844266285141221

```
In [28]: 1 plt.figure(figsize = (10, 10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha$
3 #plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge; $\alpha = 100$')
4 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
5 plt.xticks(rotation = 90)
6 plt.legend()
7 plt.show()
```




```
In [29]: 1 #Lasso regression model
2 print("\nLasso Model: \n")
3 lasso = Lasso(alpha = 10)
4 lasso.fit(X_train,y_train)
5 train_score_ls =lasso.score(X_train,y_train)
6 test_score_ls =lasso.score(X_test,y_test)
7 print("The train score for ls model is {}".format(train_score_ls))
8 print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.0
The test score for ls model is -0.0042092253233847465

```
In [30]: 1 pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

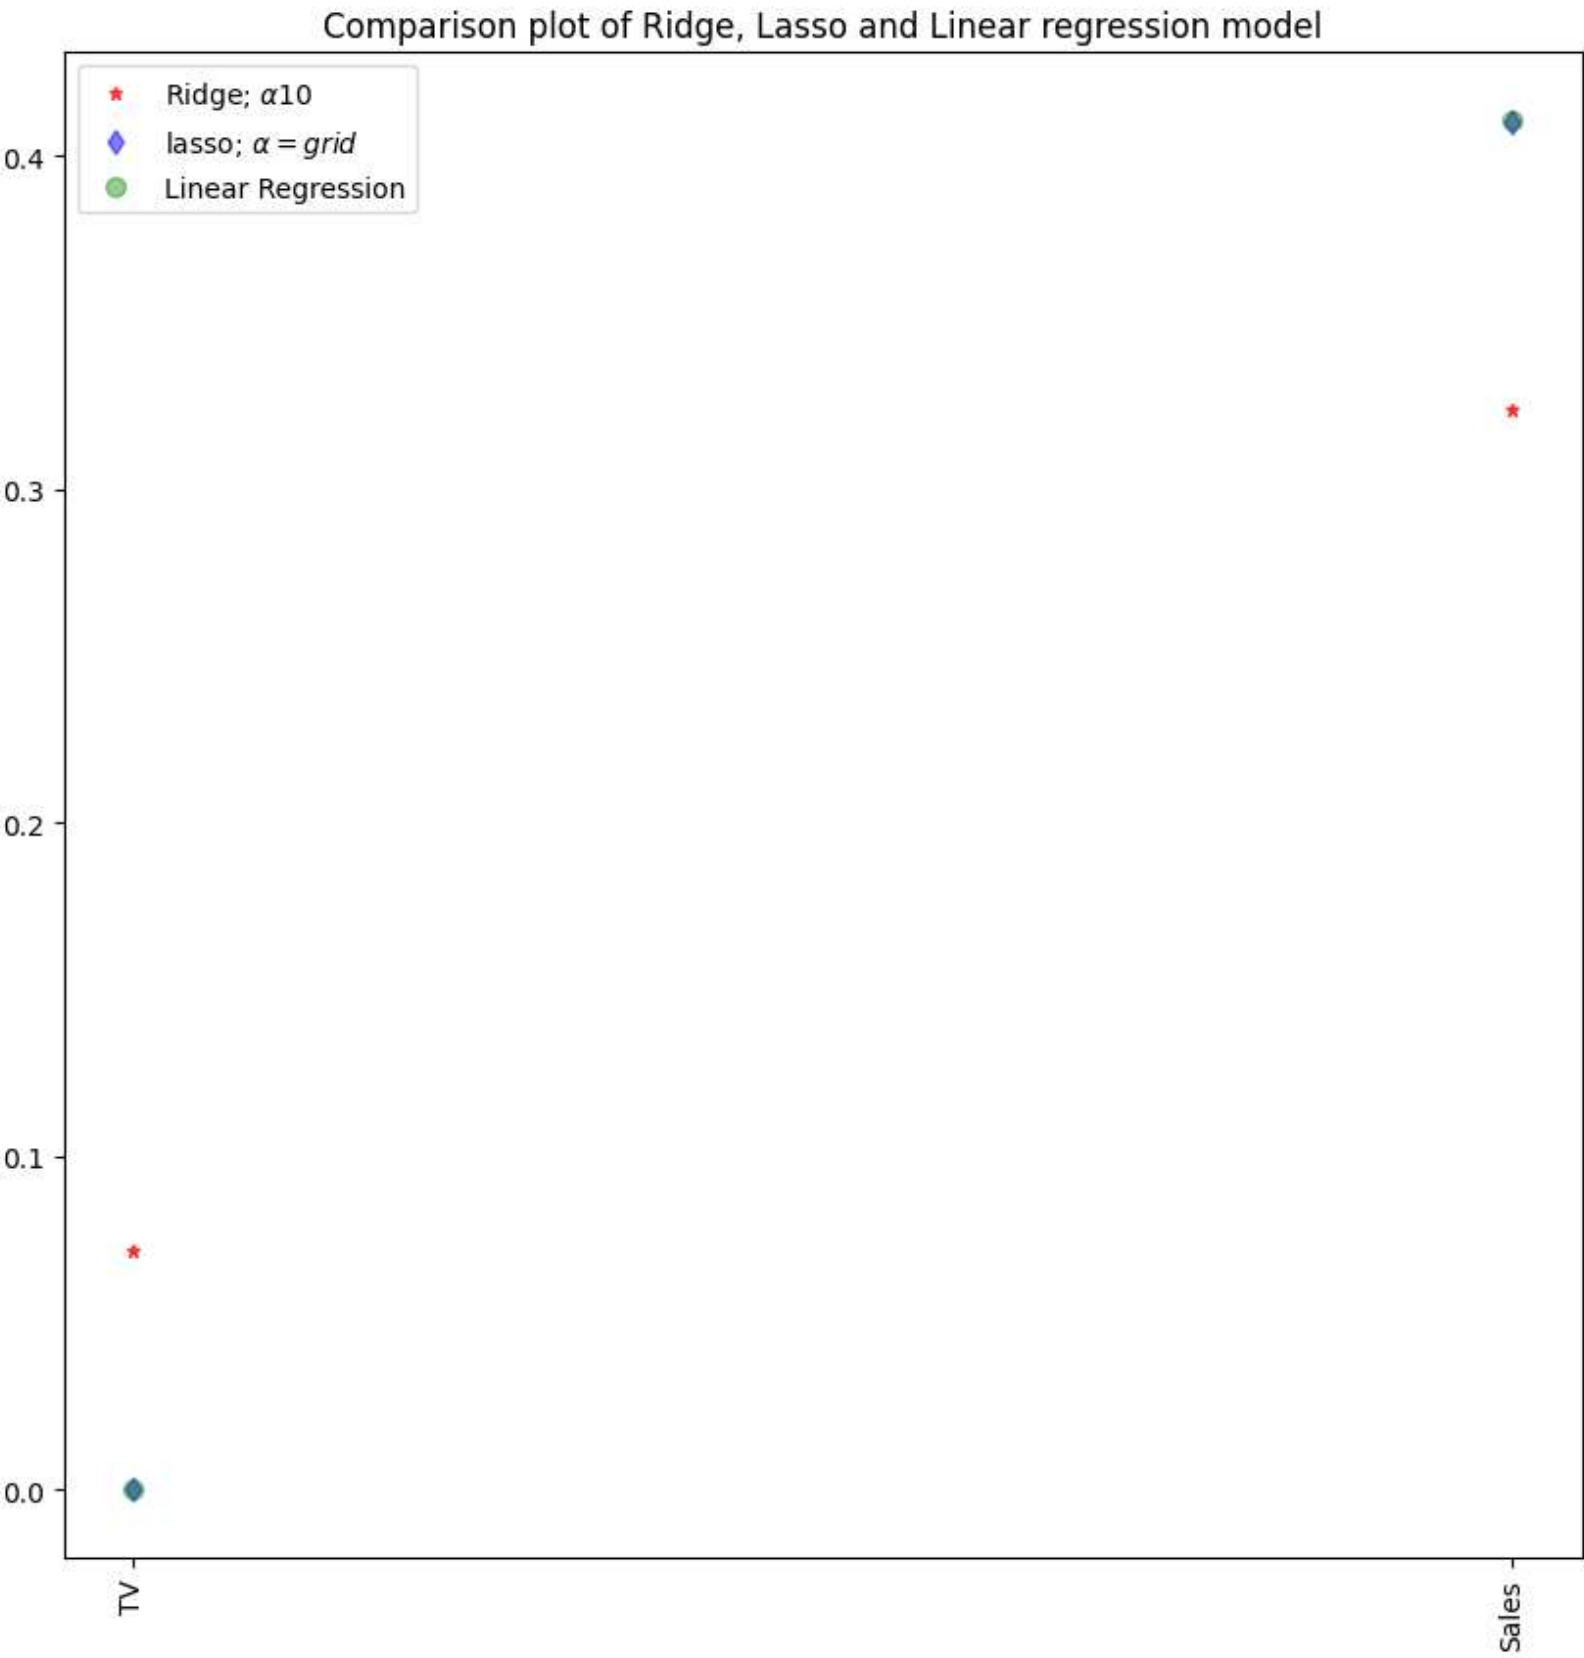
Out[30]: <Axes: >



```
In [31]: 1 #Using the linear CV model
2 from sklearn.linear_model import LassoCV
3 #Lasso Cross validation
4 lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
5
6 #score
7 print(lasso_cv.score(X_train, y_train))
8 print(lasso_cv.score(X_test, y_test))
```

0.9999999343798134
0.9999999152638072

```
In [32]: 1 #plot size
2
3 plt.figure(figsize = (10, 10))
4 #add plot for ridge regression
5 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha1$')
6 #add plot for lasso regression
7 plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso; $\alpha = grid$')
8 #add plot for linear model
9 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
10 #rotate axis
11 plt.xticks(rotation = 90)
12 plt.legend()
13 plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
14 plt.show()
15
```



```
In [33]: 1 #Using the linear CV model
2 from sklearn.linear_model import RidgeCV
3 #Ridge Cross validation
4 ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
5 #score
6 print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
7 print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.99999999997627
The train score for ridge model is 0.999999999962467

ELASTICNET

```
In [42]: 1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(X,y)
4 print(regr.coef_)
5 print(regr.intercept_)
```

[0.00417976 0.]
2.026383919311004

```
In [44]: 1 y_pred_elastic=regr.predict(X_train)
```

```
In [45]: 1 mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
2 print("Mean squared error on test set",mean_squared_error)
```

Mean squared error on test set 0.5538818050142158

vehicles dataset

```
In [27]: 1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn import preprocessing,svm
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
```

```
In [28]: 1 d=pd.read_csv(r"C:\Users\91949\Downloads\fiat500_VehicleSelection_Dataset.csv")
2 d
```

Out[28]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price	
	0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
	1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
	2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
	3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
	4	5	pop	73	3074	106880	1	41.903221	12.495650	5700

	1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
	1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
	1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
	1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
	1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

```
In [29]: 1 d=d[['engine_power','age_in_days']]
2 d.columns=['ep','aid']
```

```
In [30]: 1 d.describe()
```

Out[30]:

	ep	aid
count	1538.000000	1538.000000
mean	51.904421	1650.980494
std	3.988023	1289.522278
min	51.000000	366.000000
25%	51.000000	670.000000
50%	51.000000	1035.000000
75%	51.000000	2616.000000
max	77.000000	4658.000000

```
In [31]: 1 d.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
Column Non-Null Count Dtype
--- -
0 ep 1538 non-null int64
1 aid 1538 non-null int64
dtypes: int64(2)
memory usage: 24.2 KB

```
In [32]: 1 d.fillna(method='ffill',inplace=True)
```

C:\Users\91949\AppData\Local\Temp\ipykernel_17204\2624214790.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
d.fillna(method='ffill',inplace=True)

```
In [33]: 1 x=np.array(d['ep']).reshape(-1,1)  
2 y=np.array(d['aid']).reshape(-1,1)
```


```
In [34]: 1 d.dropna(inplace=True)
```

C:\Users\91949\AppData\Local\Temp\ipykernel_17204\1307611603.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
d.dropna(inplace=True)

```
In [35]: 1 sns.lmplot(x = "ep", y = "aid", data = d, order = 2, ci = None)
```


Out[35]: <seaborn.axisgrid.FacetGrid at 0x240065c3510>



```
In [36]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
2 regr=LinearRegression()  
3 regr.fit(x_train,y_train)  
4 print(regr.score(x_test,y_test))
```


0.08925408916467337


```
In [37]: 1 y_pred=regr.predict(x_test)
2 plt.scatter(x_test,y_test,color='b')
3 plt.plot(x_test,y_pred,color='k')
4 plt.show()
```




```
In [38]: 1 df100=d[:][:100]
2 sns.lmplot(x='ep',y='aid',data=df100,order=1,ci=None)
```

Out[38]: <seaborn.axisgrid.FacetGrid at 0x24007ee9750>




```
In [39]: 1 df100.fillna(method='ffill',inplace=True)
2 X=np.array(df100['ep']).reshape(-1,1)
3 y=np.array(df100['aid']).reshape(-1,1)
4 df100.dropna(inplace=True)
5 X_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
6 regr=LinearRegression()
7 regr.fit(X_train,y_train)
8 print(regr.score(x_test,y_test))
9 print("Regression: ",regr.score(x_test,y_test))
10 y_pred=regr.predict(x_test)
11 plt.scatter(x_test,y_test,color='g')
12 plt.plot(x_test,y_pred,color='k')
13 plt.show()
```

0.08906920769617288
Regression: 0.08906920769617288




```
In [40]: 1 plt.figure(figsize=(10,10))
2 plt.plot(feature,ridge.coef_,alpha=0.5,marker='*',markersize=5,linestyle='None',color='red',label=r'Ridge; $\alpha10$')
3 plt.plot(feature,ridge.coef_,alpha=0.6,marker='d',markersize=6,linestyle='None',color='blue',label=r'Ridge; $\alpha = g$')
4 plt.plot(feature,lr.coef_,alpha=0.7,marker='o',markersize=7,color='green',linestyle='None',label='Linear Regression')
5 plt.xticks(rotation=90)
6 plt.title("Comparision between Ridge and linear regression")
7 plt.legend()
8 plt.show()
```



```
In [41]: 1 from sklearn.linear_model import LinearRegression
2 from sklearn.model_selection import train_test_split
3 feature=d.columns[0:3]
4 target=d.columns[-1]
5 x=d[feature].values
6 y=d[target].values
7 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25)
8 lr = LinearRegression()
9 lr.fit(x_train,y_train)
10 print(lr.score(x_test,y_test))
11 print(lr.score(x_train,y_train))
```

1.0
1.0

```
In [42]: 1 plt.figure(figsize=(10,10))
2 plt.plot(feature,ridge.coef_,alpha=0.5,marker='*',markersize=5,linestyle='None',color='red',label=r'Ridge; $\alpha10$')
3 plt.plot(feature,lr.coef_,alpha=0.7,marker='o',markersize=7,color='green',linestyle='None',label='Linear Regression')
4 plt.xticks(rotation=90)
5 plt.title("Comparision between Ridge and linear regression")
6 plt.legend()
7 plt.show()
```



```
In [43]: 1 from sklearn.linear_model import Ridge,RidgeCV,Lasso,LassoCV
2 ridge = Ridge(alpha=10)
3 ridge.fit(x_train,y_train)
4 train_score_ridge=ridge.score(x_train,y_train)
5 test_score_ridge=ridge.score(x_test,y_test)
6 print('\n Ridge method\n')
7 print('The score of ridge method is {}'.format(train_score_ridge))
8 print('The score of ridge method is {}'.format(test_score_ridge))
```

Ridge method

The score of ridge method is 1.0
The score of ridge method is 1.0

```
In [44]: 1 lasso = Lasso(alpha=10)
2 lasso.fit(x_train,y_train)
3 train_score_lasso=lasso.score(x_train,y_train)
4 test_score_lasso=lasso.score(x_test,y_test)
5 print('\n Lasso method\n')
6 print('The score of lasso method is {}'.format(train_score_lasso))
7 print('The score of lasso method is {}'.format(test_score_lasso))
```

Lasso method

The score of lasso method is 0.9999999999644131
The score of lasso method is 0.9999999999644129

```
In [45]: 1 lasso_cv= LassoCV(alphas=[0.2,0.03,0.004,0.0001,1,20]).fit(x_train,y_train)
2 train_score_lasso_cv=lasso_cv.score(x_train,y_train)
3 test_score_lasso_cv=lasso_cv.score(x_test,y_test)
4 print('\n LassoCV method\n')
5 print('The score of Lasso method is {}'.format(train_score_lasso_cv))
6 print('The score of Lasso method is {}'.format(test_score_lasso_cv))
```

LassoCV method

The score of Lasso method is 1.0
The score of Lasso method is 1.0

```
In [46]: 1 ridge_cv=RidgeCV(alphas=[1,2.3,0.2,0.3,0.4,0.5,0.6]).fit(x_train,y_train)
2 print("\n RidgeCV Method\n")
3 print("The score of Ridge method is {}".format(ridge_cv.score(x_train,y_train)))
4 print("The score of Ridge method is {}".format(ridge_cv.score(x_test,y_test)))
```

RidgeCV Method

The score of Ridge method is 0.999999998399647
The score of Ridge method is 0.999999998399633